

Airbnb Project – NYC

This notebook:

- Runs EDA with charts: price distribution, listings by neighbourhood, box plot of price by room type, price vs reviews, correlation heatmap.
- Saves figures and stats to an outputs/
- Includes an ML section (linear regression for price) and simple clustering.
- Provides instructions to run the Streamlit dashboard that ships with this project.

In [1]:

```
%pip install -q pandas numpy matplotlib plotly scikit-learn requests
```

Note: you may need to restart the kernel to use updated packages.

In [1]:

```
import os
import io
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import plotly.express as px
import requests

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.cluster import KMeans

# Paths and constants
DATA_URL = "https://data.insideairbnb.com/united-states/ny/new-york"
DATA_PATH = "listings.csv"
OUT_DIR = "outputs"
os.makedirs(OUT_DIR, exist_ok=True)

# Inline plots
%matplotlib inline
```

```
/Users/shivalimuthukumar/anaconda3/lib/python3.11/site-packages/pandas/core/arrays/masked.py:61: UserWarning: Pandas requires version '1.3.6' or newer of 'bottleneck' (version '1.3.5' currently installed).
    from pandas.core import (
```

Download the NYC listings.csv

In [2]:

```
def download_if_missing(url: str, path: str) -> str:
    if os.path.exists(path) and os.path.getsize(path) > 0:
        print(f"Found local {path}. Using it.")
        return path
    print("Downloading:", url)
    r = requests.get(url, timeout=120)
    r.raise_for_status()
    with open(path, "wb") as f:
        f.write(r.content)
    print("Saved to", path)
    return path

_ = download_if_missing(DATA_URL, DATA_PATH)
```

Downloading: <https://data.insideairbnb.com/united-states/ny/new-york-city/2025-08-01/visualisations/listings.csv> (<https://data.insideairbnb.com/united-states/ny/new-york-city/2025-08-01/visualisations/listings.csv>)
Saved to listings.csv

Load and clean data

In [3]:

```
def coerce_price(series: pd.Series) -> pd.Series:
    try:
        return (
            series.astype(str)
            .str.replace("$", "", regex=False)
            .str.replace(",", "", regex=False)
            .str.replace(" ", "", regex=False)
            .str.extract(r'([0-9]*\.[0-9]+)')[0]
            .astype(float)
        )
    except Exception:
        return pd.to_numeric(series, errors="coerce")

def basic_clean(df: pd.DataFrame) -> pd.DataFrame:
    df = df.drop_duplicates().copy()
    if "price" in df.columns:
        df["price"] = coerce_price(df["price"])
    # Standard numeric casting and missing handling
    for c in df.columns:
        if df[c].dtype.kind in "biufc":
            if df[c].isna().mean() < 0.95:
                df[c] = pd.to_numeric(df[c], errors="coerce")
                df[c] = df[c].fillna(df[c].median())
        else:
            df[c] = df[c].fillna("Unknown")
    # Parse date-like columns when present
```

```

for c in [col for col in df.columns if "date" in col.lower()]:
    df[c] = pd.to_datetime(df[c], errors="coerce")
return df

raw = pd.read_csv(DATA_PATH, low_memory=False)
df = basic_clean(raw)
print("Rows:", len(df), "| Columns:", len(df.columns))
df.head()

```

Rows: 36403 | Columns: 18

Out[3]:

	id	name	host_id	host_name	neighbourhood_group	neighbourhood	latitude	longitude
0	2539	Superfast Wi-Fi. Clean & quiet home by the park	2787	John	Brooklyn	Kensington	40.64529	-74.02445
1	2595	Skylit Studio Oasis Midtown Manhattan	2845	Jennifer	Manhattan	Midtown	40.75356	-74.00890
2	6848	Only 2 stops to Manhattan studio	15991	Allen	Brooklyn	Williamsburg	40.70935	-74.01290
3	6872	Uptown Sanctuary w/ Private Bath (Month to Month)	16104	Kahshanna	Manhattan	East Harlem	40.80107	-74.01290
4	6990	UES Beautiful Blue Room	16800	Cynthia	Manhattan	East Harlem	40.78778	-74.00590

Save cleaned CSV and summary statistics

In [4]:

```
cleaned_csv_path = os.path.join(OUT_DIR, "cleaned_listings.csv")
df.to_csv(cleaned_csv_path, index=False)
print("Saved:", cleaned_csv_path)

summary = df.select_dtypes(include=[np.number]).describe()
summary_path = os.path.join(OUT_DIR, "summary_stats.csv")
summary.to_csv(summary_path)
print("Saved:", summary_path)

summary.head()
```

Saved: outputs/cleaned_listings.csv

Saved: outputs/summary_stats.csv

Out [4]:

	id	host_id	latitude	longitude	price	minimum_night
count	3.640300e+04	3.640300e+04	36403.000000	36403.000000	36403.000000	36403.000000
mean	4.525526e+17	1.737617e+08	40.728443	-73.947333	324.117792	28.617
std	5.320711e+17	1.922877e+08	0.056336	0.055033	2431.257283	29.288
min	2.539000e+03	1.678000e+03	40.500366	-74.251907	3.000000	1.000
25%	2.147279e+07	1.777536e+07	40.688320	-73.983517	130.000000	30.000

KPIs

In [6]:

```
def safe_col(df, candidates, default=None):
    for c in candidates:
        if c in df.columns:
            return c
    return default

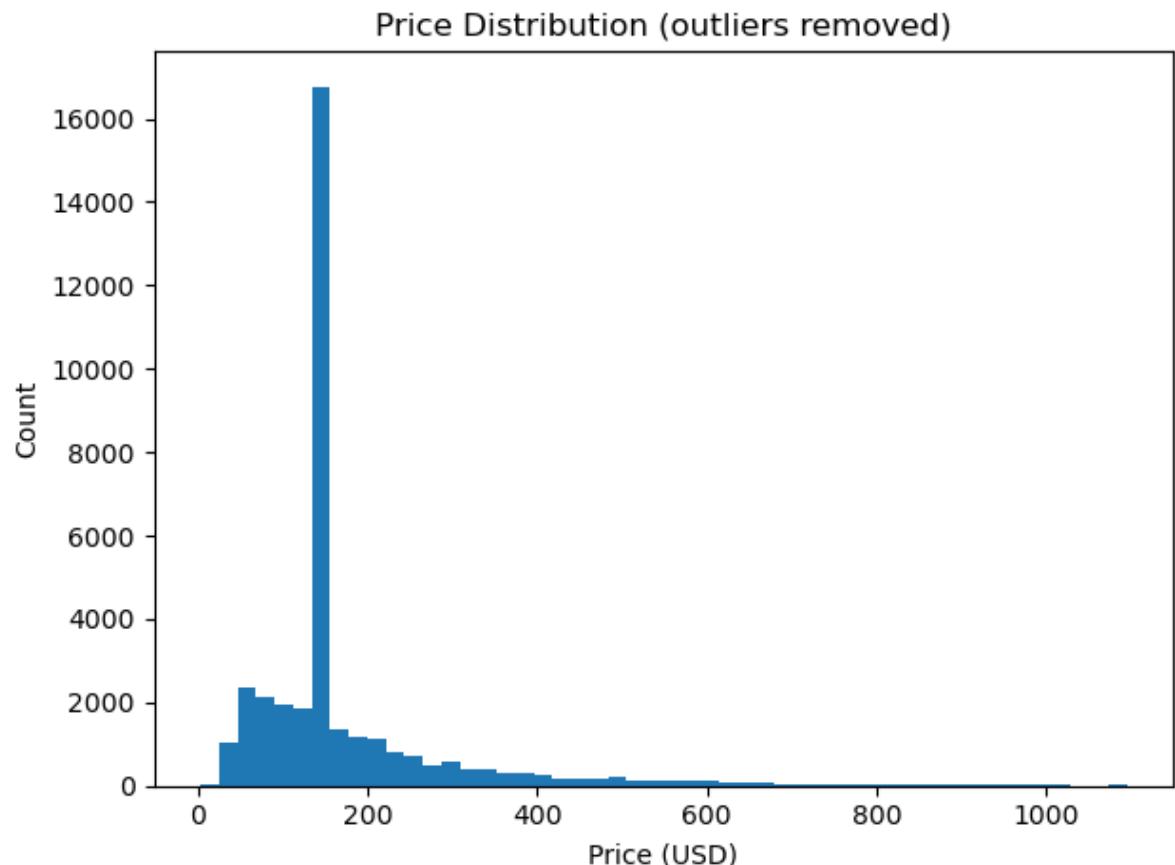
price_col = safe_col(df, ["price"])
neigh_col = safe_col(df, ["neighbourhood_cleansed", "neighbourhood"])
room_col = safe_col(df, ["room_type"])
reviews_col = safe_col(df, ["number_of_reviews"])
avail_col = safe_col(df, ["availability_365"])

kpis = {
    "Total Listings": len(df),
    "Average Price": float(df[price_col].mean()) if price_col else None,
    "Average Reviews": float(df[reviews_col].mean()) if reviews_col else None,
    "Average Availability 365": float(df[avail_col].mean()) if avail_col else None
}
kpis
```

Out[6]: {'Total Listings': 36403, 'Average Price': 324.11779248962995, 'Average Reviews': 26.85105623162926, 'Average Availability 365': 161.65516578303985}

Price distribution

```
In [8]: if price_col:  
    # Remove extreme outliers (e.g., top 1%)  
    price_series = df[price_col].dropna()  
    cutoff = price_series.quantile(0.99)  
    filtered_prices = price_series[price_series <= cutoff]  
  
    plt.figure()  
    filtered_prices.plot(kind="hist", bins=50)  
    plt.xlabel("Price (USD)")  
    plt.ylabel("Count")  
    plt.title("Price Distribution (outliers removed)")  
    plt.tight_layout()  
    out = os.path.join(OUT_DIR, "price_hist_filtered.png")  
    plt.savefig(out, dpi=150)  
    plt.show()  
    print("Saved:", out)  
else:  
    print("No price column found.")
```



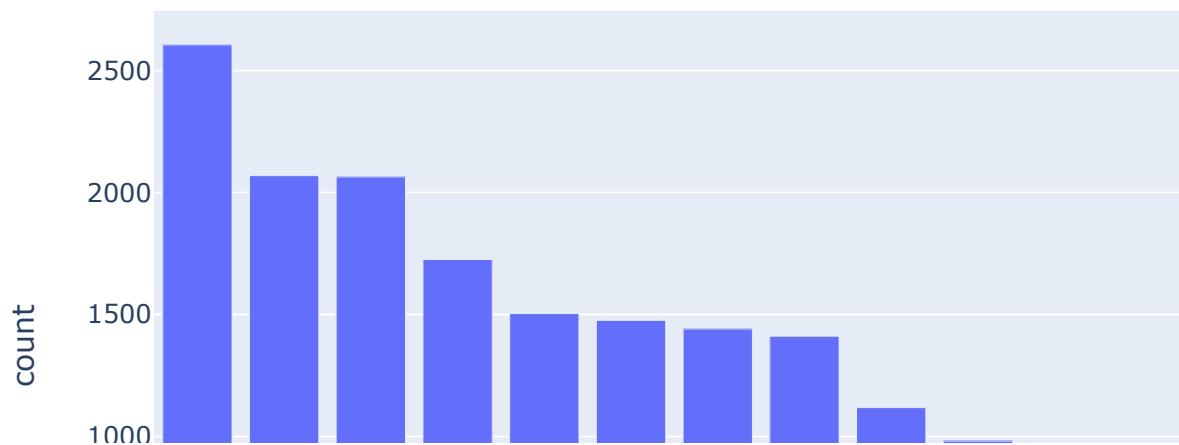
Saved: outputs/price_hist_filtered.png

Listings by neighbourhood

In [9]:

```
if neigh_col:  
    topN = 20  
    vc = df[neigh_col].value_counts().head(topN).reset_index()  
    vc.columns = [neigh_col, "count"]  
    fig = px.bar(vc, x=neigh_col, y="count", title=f"Top {topN} nei  
    fig.show()  
    out_html = os.path.join(OUT_DIR, "listings_by_neighbourhood.htm  
    fig.write_html(out_html, include_plotlyjs="cdn")  
    print("Saved:", out_html)  
else:  
    print("No neighbourhood column found.")
```

Top 20 neighbourhoods by listing count



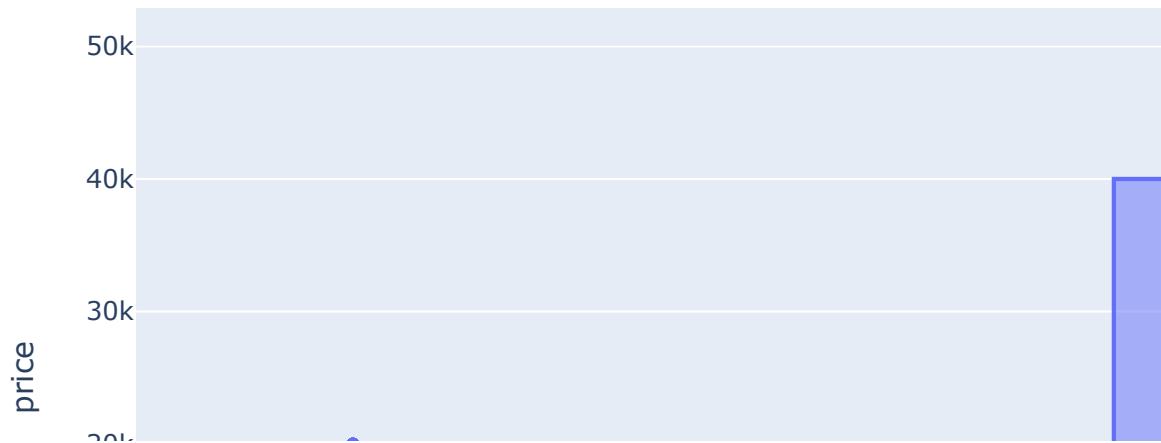
Saved: outputs/listings_by_neighbourhood.html

Price by room type

In [10]:

```
if price_col and room_col:  
    fig = px.box(df[[room_col, price_col]].dropna(), x=room_col, y=price_col)  
    fig.show()  
    out_html = os.path.join(OUT_DIR, "price_by_room_type.html")  
    fig.write_html(out_html, include_plotlyjs="cdn")  
    print("Saved:", out_html)  
else:  
    print("Room type or price column missing.")
```

Price by room type

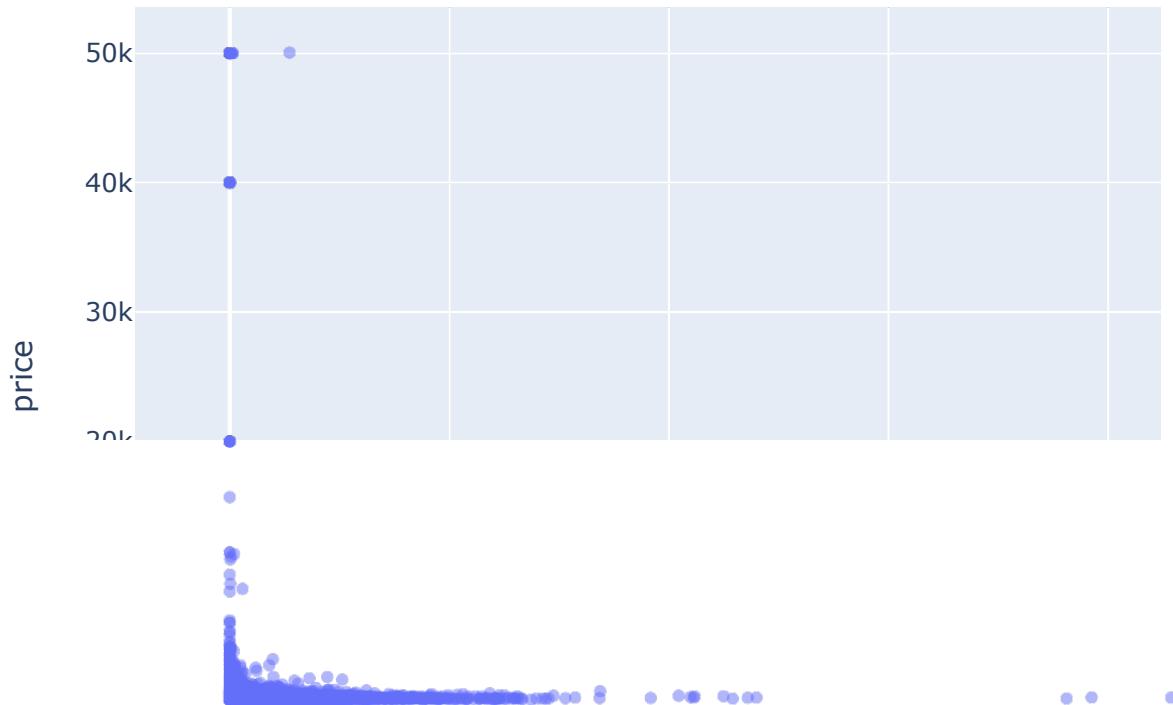


Saved: outputs/price_by_room_type.html

Price vs number of reviews

```
In [12]: if price_col and reviews_col:  
    fig = px.scatter(  
        df,  
        x=reviews_col,  
        y=price_col,  
        opacity=0.5,  
        title="Price vs number of reviews"  
    )  
    fig.show()  
    out_html = os.path.join(OUT_DIR, "price_vs_reviews.html")  
    fig.write_html(out_html, include_plotlyjs="cdn")  
    print("Saved:", out_html)
```

Price vs number of reviews



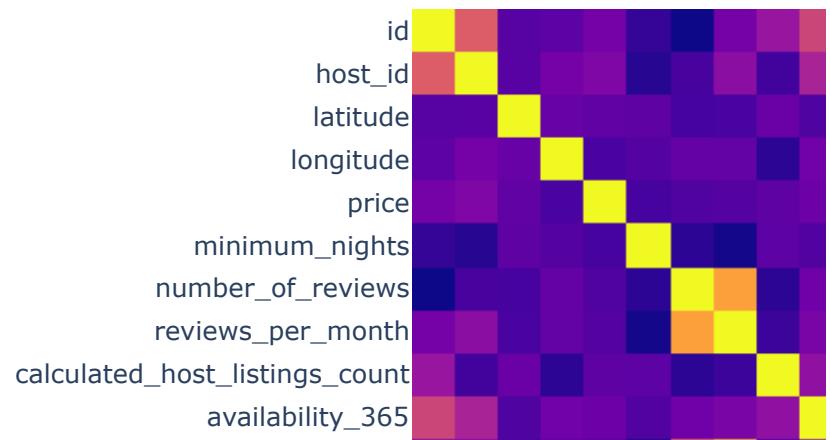
Saved: outputs/price_vs_reviews.html

Correlation heatmap

In [13]:

```
num_df = df.select_dtypes(include=[np.number])
if not num_df.empty:
    corr = num_df.corr(numeric_only=True)
    fig = px.imshow(corr, title="Correlation Heatmap")
    fig.show()
    out_html = os.path.join(OUT_DIR, "correlation_heatmap.html")
    fig.write_html(out_html, include_plotlyjs="cdn")
    print("Saved:", out_html)
else:
    print("No numeric columns for correlation.")
```

Correlation Heatmap



Saved: outputs/correlation_heatmap.html

Linear regression for price

In [16]:

```

if price_col:
    feature_candidates = [c for c in ["minimum_nights", "maximum_nights"]]
    if feature_candidates:
        X = df[feature_candidates].fillna(0.0)
        y = df[price_col].fillna(df[price_col].median())
        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
        model = LinearRegression().fit(X_train, y_train)
        r2 = model.score(X_test, y_test)
        print("R^2 on hold-out:", round(r2, 3))
        coef = pd.DataFrame({"feature": feature_candidates, "coef": model.coef_})
        display(coef)
    else:
        print("Not enough numeric features for regression.")
else:
    print("Price column is required for regression.")

```

R^2 on hold-out: 0.008

	feature	coef
1	availability_365	0.939912
3	calculated_host_listings_count	0.086410
2	number_of_reviews	-1.123670
0	minimum_nights	-4.246441

Simple K-Means clustering

In [19]:

```

num_df = df.select_dtypes(include=[np.number])
if not num_df.empty:
    k = 3
    km = KMeans(n_clusters=k, n_init="auto", random_state=42)
    fit_df = num_df.fillna(num_df.median(numeric_only=True))
    labels = km.fit_predict(fit_df)
    df_clu = df.copy()
    df_clu["cluster"] = labels
    print("Cluster counts:")
    print(df_clu["cluster"].value_counts())
    display(df_clu.head(10))
else:
    print("No numeric columns available for clustering.")

import matplotlib.pyplot as plt

plt.scatter(fit_df.iloc[:, 0], fit_df.iloc[:, 1], c=labels, cmap="viridis")
plt.title("K-Means Clusters (first 2 numeric features)")
plt.xlabel(fit_df.columns[0])

```

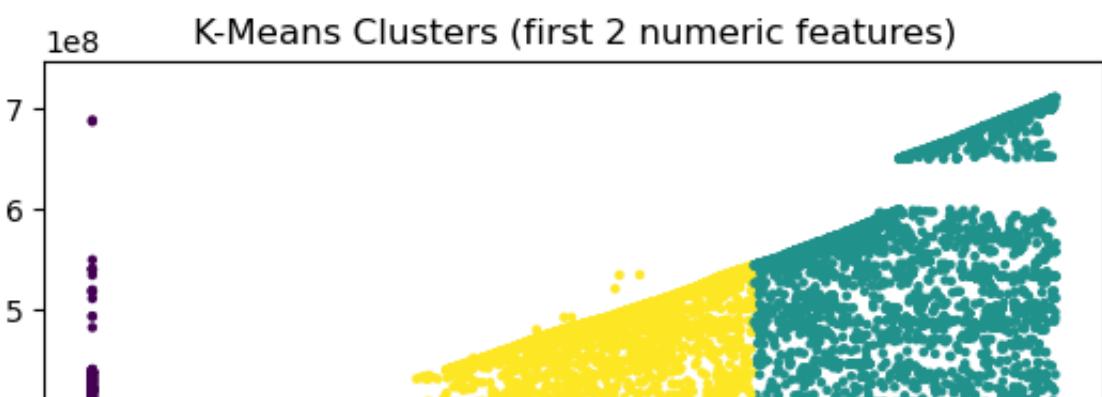
```
plt.ylabel(fit_df.columns[1])
plt.show()
```

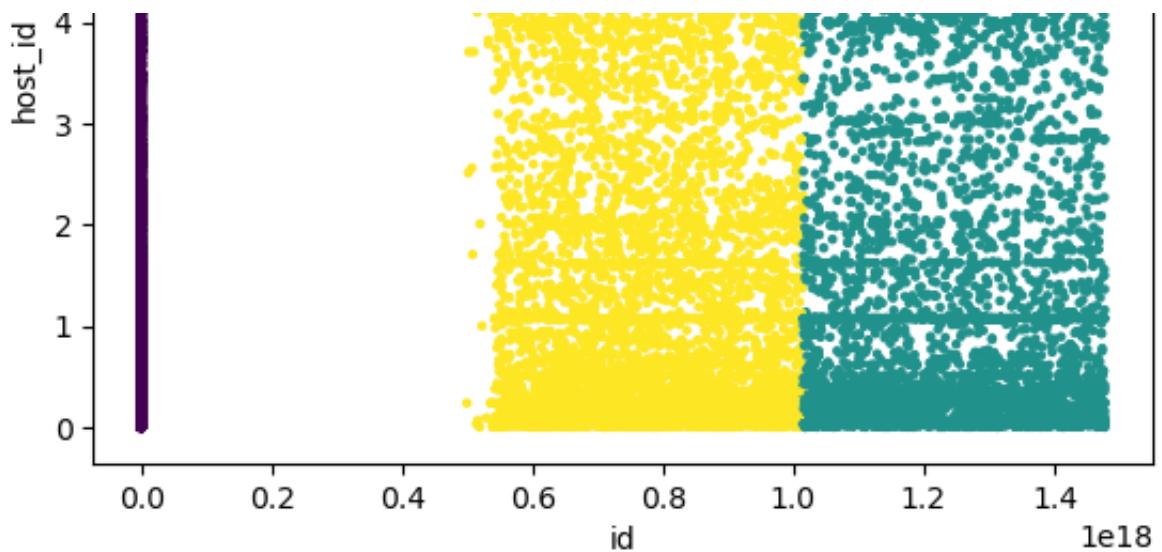
Cluster counts:

cluster	count
0	20100
2	8325
1	7978

Name: count, dtype: int64

	id	name	host_id	host_name	neighbourhood_group	neighbourhood	latitude	longitude
0	2539	Superfast Wi-Fi. Clean & quiet home by the park	2787	John	Brooklyn	Kensington	40.6532	-74.0475
1	2595	Skylit Studio Oasis Midtown Manhattan	2845	Jennifer	Manhattan	Midtown	40.7128	-74.0128
2	6848	Only 2 stops to Manhattan studio	15991	Allen	Brooklyn	Williamsburg	40.6731	-74.0236
3	6872	Uptown Sanctuary w/ Private Bath (Month to Month)	16104	Kahshanna	Manhattan	East Harlem	40.7894	-74.0060
4	6990	UES Beautiful Blue Room	16800	Cynthia	Manhattan	East Harlem	40.7894	-74.0060
5	7064	Amazing location! Wburg. Large, bright & tranquil	17297	Joelle	Brooklyn	Williamsburg	40.6731	-74.0236
6	7097	Perfect for Your Parents, With Garden & Patio	17571	Jane	Brooklyn	Fort Greene	40.6731	-74.0236
7	7801	Sunny Williamsburg Loft with Sauna	21207	Chaya	Brooklyn	Williamsburg	40.6731	-74.0236
8	8490	Maison des Sirenes1,bohemian, luminous apartment	25183	Nathalie	Brooklyn	Bedford-Stuyvesant	40.6731	-74.0236
9	9357	Midtown Pied-a-terre	30193	Tommi Laurelle	Manhattan	Hell's Kitchen	40.7894	-74.0060





Export example filtered CSV

In [20]:

```
room_col = room_col or "room_type" if "room_type" in df.columns else None
if room_col and price_col:
    filtered = df[(df[room_col] == "Entire home/apt") & (df[price_col] <= max_price)]
    filtered_path = os.path.join(OUT_DIR, "filtered_listings.csv")
    filtered.to_csv(filtered_path, index=False)
    print("Saved:", filtered_path, "| Rows:", len(filtered))
else:
    print("Cannot create filtered CSV because room_type or price is missing")
```

Saved: outputs/filtered_listings.csv | Rows: 16288

How to run the Streamlit dashboard

The Streamlit dashboard is provided in `app.py`. It can auto-download the same NYC dataset if `listings.csv` is missing.

Run these commands in terminal from the project folder:

```
pip install -r requirements.txt
python fetch_nyc_data.py
streamlit run app.py
```

Then open the local URL that Streamlit prints (usually <http://localhost:8501> (<http://localhost:8501>)). Use the sidebar to filter and the buttons to export `cleaned_listings.csv` and `filtered_listings.csv`.

In []: